

Privacy-Preserving Financial Fraud Detection Using Federated Learning on Non-IID Data

Your Name

January 8, 2026

Abstract

Financial fraud detection is a critical challenge for banking institutions, where the sophistication of fraudulent activities is constantly increasing. Traditional centralized machine learning approaches require aggregating sensitive transaction data into a central repository, which raises significant privacy concerns and violates strict data protection regulations. This thesis proposes a Federated Learning (FL) framework utilizing Neural Networks to enabling collaborative fraud detection across multiple banking institutions without sharing raw data. By employing the Federated Averaging (FedAvg) algorithm and a standardized feature engineering schema, the proposed system ensures data privacy while leveraging the collective intelligence of the network. We address key challenges such as non-IID data distribution and communication efficiency. The proposed methodology offers a robust, scalable, and privacy-preserving solution for modern financial security infrastructures.

Contents

1	Introduction	5
1.1	Research Topic	5
1.1.1	Research Motivation	5
1.1.2	Scientific Novelty	5
1.1.3	Practical Relevance	5
1.2	Detailed Research Proposal	6
1.2.1	Problem Statement	6
1.2.2	Research Gap	6
1.2.3	Research Objectives	6
1.2.4	Proposed Methodology	6
1.2.5	Contributions	7
1.2.6	Privacy and Ethical Considerations	7
2	Literature Review	8
2.1	Federated Learning Theory	8
2.2	Challenges in Federated Learning	8
2.2.1	Non-IID Data	8
2.2.2	Privacy and Security	8
2.3	Financial Fraud Detection and Federated Learning	9
2.4	Research Gaps	9
3	Methodology	10
3.1	Research Questions and Hypotheses	10
3.1.1	Research Questions (RQs)	10
3.1.2	Hypotheses	11
3.2	Theory and Model Selection	11
3.2.1	Choice of Model: Neural Networks	11
3.2.2	Federated Averaging (FedAvg)	11
3.3	Data and Feature Engineering	12
3.3.1	Data Sources	12
3.3.2	Feature Standardization Rule	12
3.3.3	Handling Missing Data	12
3.4	System Workflow	12
3.5	Data Harmonization & Feature Engineering	13
3.5.1	Data Pipeline Overview	13
3.5.2	Feature Engineering Details	14
3.5.3	Handling Schema Mismatch	14
3.6	System Workflow	15

3.6.1	Phase 1: Initialization	15
3.6.2	Phase 2: Distribution & Local Training	16
3.6.3	Phase 3: Aggregation & Termination	16
4	Evaluation and Security Discussion	18
4.1	Evaluation Metrics	18
4.2	Analysis of Non-IID Data Effects	18
4.3	Security and Privacy Discussion	18
4.3.1	Privacy Leakage Risks	18
4.3.2	Defensive Mechanisms	19
4.4	Comparison with Centralized Learning	19
5	Conclusion and Future Work	20
5.1	Conclusion	20
5.1.1	Summary of Findings	20
5.1.2	Final Remarks	20
5.2	Contributions	20
5.3	Future Research Directions	21

List of Figures

3.1	High-Level System Architecture: Hub-and-Spoke Topology. The Central Aggregator coordinates training without access to private banking data.	13
3.2	Data Harmonization Pipeline: Transforming heterogeneous raw data into the Standardized Feature Vector (X).	14
3.3	Neural Network Architecture designed for Federated Fraud Detection.	15
3.4	Sequence Diagram of a single Federated Learning training round.	17

List of Tables

3.1 Standardized Feature Schema for FL	12
--	----

Chapter 1

Introduction

1.1 Research Topic

1.1.1 Research Motivation

In the rapidly evolving digital finance landscape, the volume of online transactions has surged, accompanied by increasingly sophisticated fraudulent activities. Banking institutions are under immense pressure to detect and prevent financial fraud to protect assets and maintain reputation. Traditional fraud detection systems often operate in silos, where each bank relies solely on its internal data. However, fraud patterns are dynamic and often cross institutional boundaries. While sharing data between banks could significantly enhance detection capabilities, strict data privacy regulations (such as GDPR, CCPA) and competitive confidentiality prevent the direct exchange of raw transaction data. This creates a dilemma: banks need to collaborate to improve security but are legally and practically restricted from doing so.

1.1.2 Scientific Novelty

This research proposes a Federated Learning (FL) framework specifically designed for the banking sector. Unlike standard FL applications in mobile keyboard prediction, financial fraud detection involves highly imbalanced and non-Independent and Identically Distributed (non-IID) data. This thesis addresses the specific challenge of applying Neural Networks within a Federated Averaging (FedAvg) context to learn global fraud patterns from local, heterogeneous banking data without raw data transmission.

1.1.3 Practical Relevance

The proposed solution offers a practical architecture that allows banks to benefit from "community knowledge" of fraud vectors without exposing sensitive customer information. This aligns with the strategic needs of modern Fintech infrastructure, balancing high-accuracy security with strict regulatory compliance.

1.2 Detailed Research Proposal

1.2.1 Problem Statement

Financial institutions face a dual challenge: the need for high-accuracy fraud detection models that require vast amounts of diverse data, and the imperative to comply with data privacy laws that prohibit data sharing. Centralized Machine Learning (ML), which aggregates data into a single server, is no longer feasible due to these privacy and security risks. Consequently, models trained on isolated data (data islands) suffer from poor generalization, especially against new or rare fraud types that have not yet appeared in a specific bank's local dataset but may be prevalent elsewhere.

1.2.2 Research Gap

Existing literature extensively covers FL for edge computing (e.g., smartphones) but has limited in-depth exploration of FL in cross-silo banking environments where data is highly skewed (imbalanced classes) and follows non-IID distributions. Furthermore, while tree-based models (e.g., Random Forest) are standard in centralized fraud detection, their adaptation to FL is complex and often less efficient than gradient-based methods; however, the explicit justification and optimization of Neural Networks for this specific tabular domain in a federated setting require further investigation.

1.2.3 Research Objectives

The primary objective is to develop a privacy-preserving fraud detection system using Federated Learning. Specific objectives include:

1. To design a specialized Neural Network architecture suitable for tabular transaction data in a federated setting.
2. To implement a standardized feature engineering schema that ensures vector consistency across different banking entities.
3. To evaluate the performance of Federated Averaging (FedAvg) on non-IID financial data compared to local-only training.
4. To analyze the privacy implications and ensuring that model updates do not leak sensitive information.

1.2.4 Proposed Methodology

The research employs a horizontal Federated Learning approach. A central server coordinates the training process by initializing a global Neural Network model. Participating banks (clients) download this model, train it locally on their private, labeled transaction data (including chargebacks and verified fraud cases), and compute model updates (gradients/weights). These updates are sent back to the server, which aggregates them using the FedAvg algorithm to update the global model. This cycle repeats until convergence.

1.2.5 Contributions

Theoretical Contribution: This thesis contributes to the understanding of FL convergence on non-IID tabular data, specifically in the context of binary classification with extreme class imbalance. **Practical Contribution:** A deployable framework for inter-bank collaboration is provided, including specific protocols for data standardization and model aggregation that respect privacy boundaries.

1.2.6 Privacy and Ethical Considerations

The system adheres to "Privacy by Design" principles. Raw data never leaves the local premise. The research also considers the risk of model inversion attacks and discusses the integration of Differential Privacy and Secure Aggregation to mitigate potential inference risks from shared gradients.

Chapter 2

Literature Review

2.1 Federated Learning Theory

The concept of Federated Learning was formally introduced by [6], who proposed the Federated Averaging (FedAvg) algorithm. This seminal work demonstrated that it is possible to train deep neural networks on decentralized data without determining the raw data to a central location. The authors addressed the key challenge of communication efficiency by performing multiple steps of local stochastic gradient descent (SGD) on each client before averaging the weights at the server.

Building on this, [4] further explored optimized strategies for distributed machine learning, emphasizing the reduction of uplink communication costs through structured updates and sketched updates. [7] provided a comprehensive categorization of FL into Horizontal FL, Vertical FL, and Federated Transfer Learning, establishing the standard taxonomy used in the field. They highlighted the applicability of FL in financial risk management, where institutions share overlapping user samples but different feature spaces (Vertical FL) or overlapping feature spaces but different users (Horizontal FL), the latter being the focus of this research for inter-bank collaboration.

2.2 Challenges in Federated Learning

2.2.1 Non-IID Data

One of the most significant challenges in FL is the statistical heterogeneity of data, often referred to as non-IID (Independent and Identically Distributed) data. [5] analyzed the convergence of FedAvg in heterogeneous networks and proposed a proximal term to the local objective function (FedProx) to improve stability. In the context of banking, fraud patterns may vary significantly between institutions (e.g., a retail bank vs. an investment bank), making the non-IID assumption critical for robust model design.

2.2.2 Privacy and Security

While FL offers privacy improvements by keeping raw data local, gradient leakage remains a concern. [1] introduced a practical Secure Aggregation protocol using cryptographic techniques to ensure that the server can only inspect the aggregate of the updates, not the individual contributions, thereby preventing the reconstruction of in-

dividual user data. [3] provide an extensive survey of these privacy mechanisms and the trade-offs between privacy (e.g., Differential Privacy) and model utility.

2.3 Financial Fraud Detection and Federated Learning

Traditional fraud detection relies heavily on centralized data mining techniques. [2] discusses the challenges of credit card fraud detection, including concept drift and class imbalance, in a realistic centralized setting. In the federated domain, recent studies have begun to apply FL to this problem. However, the literature often focuses on generic credit scoring or simplified fraud scenarios. The application of FL specifically to cross-institution transactional fraud detection, dealing with the intricacies of tabular data and extreme class imbalance without sharing user identifiers, remains an active area of research.

2.4 Research Gaps

Based on the review of existing literature, several key gaps are identified:

1. **Limited Real-World Banking Deployments:** Most FL research [6, 5] focuses on benchmarks like MNIST or CIFAR-10, or mobile device data (e.g., next-word prediction). There is a scarcity of detailed methodologies for deploying FL in the specific regulatory and technical infrastructure of core banking systems.
2. **Non-IID Data Handling in Fraud Detection:** While [5] address non-IID data theoretically, the specific impact of heterogeneous fraud distributions (e.g., one bank seeing a specific attack vector while others do not) on the global model's false positive rate requires specific investigation using tabular financial data.
3. **Model Preference for Tabular Data:** In centralized fraud detection, tree-based models (Random Forest, GBM) are dominant [2]. However, these are difficult to adapt to FL due to the challenge of averaging discrete tree structures. There is a need to rigorously justify and optimize Neural Networks as the primary alternative for FL-based fraud detection to ensure they match the performance expectations of financial institutions.
4. **Standardization in Cross-Silo FL:** The literature often assumes compatible feature spaces. In a practical inter-bank scenario, defining a rigid, privacy-preserving feature engineering schema that aligns disparate raw data schemas remains a practical gap that this thesis aims to address.

Chapter 3

Methodology

3.1 Research Questions and Hypotheses

3.1.1 Research Questions (RQs)

To address the identified gaps, this thesis poses the following research questions:

1. **RQ1:** How does the performance of a Neural Network trained via Federated Learning compare to models trained on isolated local data in terms of fraud detection accuracy?
2. **RQ2:** Can the Federated Averaging (FedAvg) algorithm effectively converge when participating banks have highly non-IID fraud distributions?
3. **RQ3:** What is the impact of participating in FL on the False Positive Rate (FPR) for individual banks?
4. **RQ4:** Is a Neural Network architecture sufficient to replace traditional tree-based models for tabular fraud data in a federated setting?
5. **RQ5:** How does the number of participating banks affect the convergence speed of the global model?
6. **RQ6:** Can a standardized feature engineering schema effectively bridge the gap between heterogeneous raw banking databases?
7. **RQ7:** What is the communication overhead required for synchronizing model parameters in a realistic banking network?
8. **RQ8:** Does FL improve detection of "new" fraud patterns that a specific bank has not yet encountered?
9. **RQ9:** How resilient is the proposed FL system to a participating bank dropping out during training?
10. **RQ10:** What are the minimum local training epochs key to balancing communication cost and model accuracy?

3.1.2 Hypotheses

From these questions, we derive the following testable hypotheses:

1. **H1:** The global FL model will achieve a higher Area Under the Precision-Recall Curve (AUPRC) than the average AUPRC of models trained only on local data.
2. **H2:** Participating in FL will reduce the False Negative Rate for "rare" fraud types that are present in the collective dataset but sparse locally.
3. **H3:** The FedAvg algorithm will converge to a stable global model even when local fraud rates vary by an order of magnitude (non-IID).
4. **H4:** A Neural Network with appropriate embedding layers for categorical variables will perform comparably to a centralized Random Forest baseline.
5. **H5:** Increasing the number of local epochs (E) reduces the total communication rounds required for convergence.
6. **H6:** The standardized feature vector approach allows for effective model aggregation without requiring raw data alignment.

3.2 Theory and Model Selection

3.2.1 Choice of Model: Neural Networks

While tree-based models like Random Forest and XGBoost are industry standards for tabular data, they pose significant challenges in a Federated Learning context. Aggregating decision trees from different clients is non-trivial because trees partition the feature space differently based on local data distributions. Averaging tree structures or leaf values is mathematically complex and often leads to performance degradation or privacy leakage (via threshold exposure).

In contrast, **Neural Networks (NN)** are chosen for this research because:

- **Gradient-Based Optimization:** NNs are optimized using stochastic gradient descent (SGD), which naturally fits the FedAvg algorithm [6].
- **Parameter Aggregation:** The weights of a NN are real-valued matrices that can be averaged element-wise. This provides a clear mathematical foundation for the aggregation step: $w_{global} = \sum \frac{n_k}{n} w_k$.
- **Non-Linearity:** NNs can learn complex, non-linear patterns in fraud data through hidden layers and activation functions (ReLU), approximating the decision boundaries that decision trees would find.

3.2.2 Federated Averaging (FedAvg)

The core algorithm used is FedAvg. The optimization objective is:

$$\min_w F(w) = \sum_{k=1}^K \frac{n_k}{n} F_k(w) \quad (3.1)$$

where $F_k(w)$ is the local loss function at bank k . In each round t , the server distributes the current global model w_t . Each bank k performs E epochs of local training: $w_{t+1}^k \leftarrow w_t - \eta \nabla F_k(w_t)$. The server then aggregates: $w_{t+1} \leftarrow \sum_{k=1}^K \frac{n_k}{n} w_{t+1}^k$.

3.3 Data and Feature Engineering

3.3.1 Data Sources

The system relies on internal transaction databases at each bank. Labels ($y \in \{0, 1\}$) are derived from:

- Customer disputes and chargebacks.
- Investigations by fraud analyst teams.
- Confirmed rule-based triggers.

Crucially, labels are accurate ground truths derived from historical outcomes, not generated by the model itself.

3.3.2 Feature Standardization Rule

A critical constraint of this FL framework is that while *raw data* storage formats may differ across banks (e.g., 'Oracle' vs 'PostgreSQL', different column names), the **Feature Vector** input to the model must be identical in structure.

Defined Schema:

Feature	Description
<code>txn_amount</code>	Normalized transaction value
<code>txn_hour</code>	Hour of day (0-23)
<code>merchant_cat</code>	Category code (One-hot encoded)
<code>is_foreign</code>	Binary flag for international txn
<code>device_hash</code>	High-cardinality categorical (Embedding)

Table 3.1: Standardized Feature Schema for FL

3.3.3 Handling Missing Data

If a bank does not collect a specific feature (e.g., `device_hash`), the protocol mandates filling this with a sentinel value (e.g., 0 or 'UNKNOWN') rather than omitting the column, ensuring the vector dimensions remain $d \times 1$ for all clients.

3.4 System Workflow

The operational workflow for the proposed Federated Learning system is as follows:

1. **Initialization:** The central server utilizes a random seed or a pre-trained base to initialize the weights w_0 of the Neural Network.

2. **Distribution:** The server broadcasts w_0 to all participating banks K .

3. Local Training (Client-Side):

- Bank k loads w_0 into its local infrastructure.
- It trains the model on its private dataset D_k for E epochs using a local optimizer (e.g., Adam).
- Result: Local weights w_t^k .

4. **Upload:** Bank k uploads the weight update $\Delta w = w_t^k$ (or the weights themselves) to the server. This step is secured via TLS and potentially Secure Aggregation.

5. **Aggregation (Server-Side):** The server computes the weighted average of the received models: $w_{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_t^k$.

6. **Iteration:** The new global model w_{t+1} is broadcast back to banks. Steps 3-5 repeat until the loss function converges or a fixed number of rounds is reached.

7. **Deployment:** The final converged model is frozen and deployed into the production transaction processing pipeline of each bank for real-time inference.

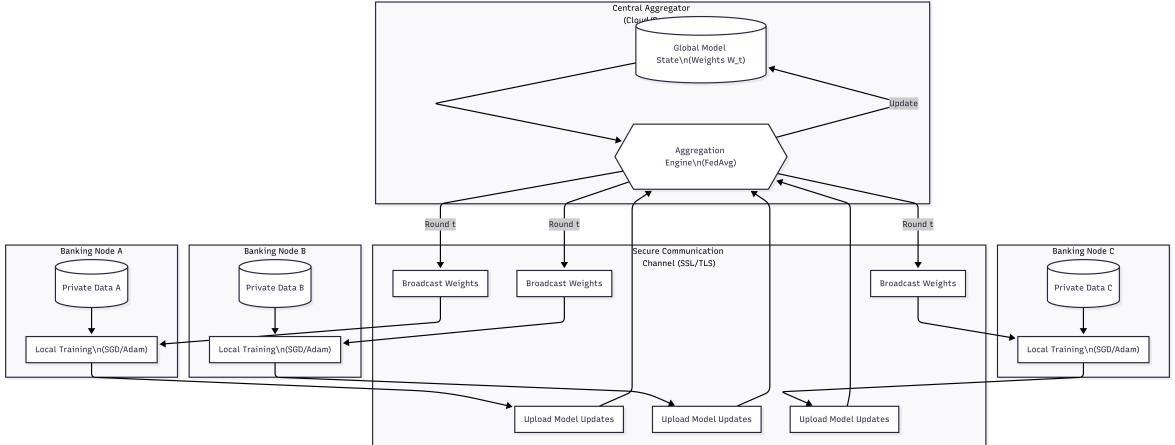


Figure 3.1: High-Level System Architecture: Hub-and-Spoke Topology. The Central Aggregator coordinates training without access to private banking data.

3.5 Data Harmonization & Feature Engineering

3.5.1 Data Pipeline Overview

Given the heterogeneity of raw data across banks, a robust data pipeline is essential. This pipeline ensures that despite differences in raw data formats or missing fields, the input to the Neural Network remains consistent and standardized.

- **Raw Data Collection:** Ingesting transaction data from various sources within the bank.

- **Schema Mapping:** Aligning incoming data to the predefined Common Schema.
- **Feature Engineering:** Deriving additional features or transforming existing ones to better represent the underlying patterns.
- **Normalization:** Scaling features to a standard range, ensuring no single feature dominates due to its scale.
- **Encoding:** Converting categorical variables into a numerical format that can be fed into the Neural Network.

3.5.2 Feature Engineering Details

To implement the Common Schema, each bank performs specific transformations on its raw data:

Feature	Type	Transformation
txn_amount	Float	Min-Max Scaler: Scales the transaction amount to [0, 1].
txn_hour	Int	One-Hot Encoder: 24 binary features for each hour.
merchant_cat	Categorical	Target Encoder: Replaces category with mean target value.
is_foreign	Binary	Leave As Is: Already in a suitable format.
device_hash	Categorical	Frequency Encoder: Replaces hash with its frequency.
dist_dev	Float	Standard Scaler: Z-score normalization of deviation from user's avg transaction amount.

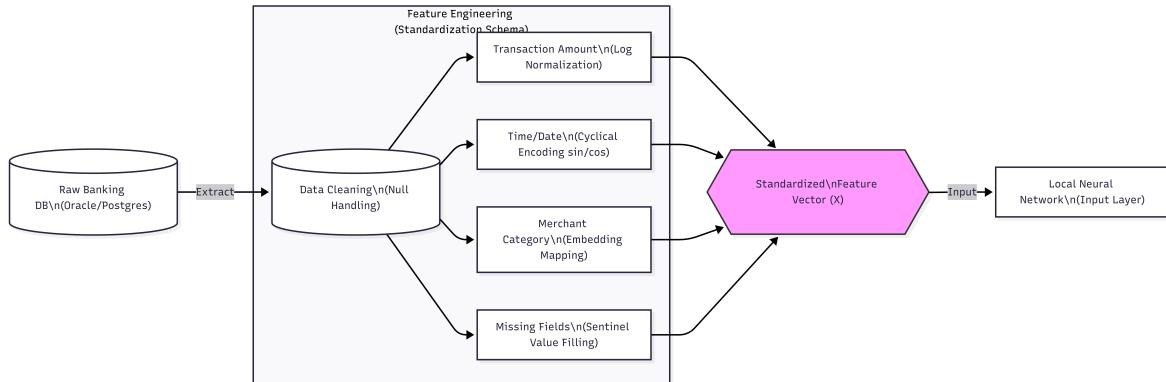


Figure 3.2: Data Harmonization Pipeline: Transforming heterogeneous raw data into the Standardized Feature Vector (X).

3.5.3 Handling Schema Mismatch

In cases where a bank's raw data schema lacks certain features present in the Common Schema:

- The missing feature is filled with a sentinel value (e.g., 0, 'UNKNOWN').
- This ensures all feature vectors have the same dimensionality, crucial for model training.

3.6 System Workflow

The operational workflow for the proposed Federated Learning system is as follows:

1. **Initialization:** The central server utilizes a random seed or a pre-trained base to initialize the weights w_0 of the Neural Network.
2. **Distribution:** The server broadcasts w_0 to all participating banks K .
3. **Local Training (Client-Side):**
 - Bank k loads w_0 into its local infrastructure.
 - It trains the model on its private dataset D_k for E epochs using a local optimizer (e.g., Adam).
 - Result: Local weights w_t^k .
4. **Upload:** Bank k uploads the weight update $\Delta w = w_t^k$ (or the weights themselves) to the server. This step is secured via TLS and potentially Secure Aggregation.
5. **Aggregation (Server-Side):** The server computes the weighted average of the received models: $w_{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_t^k$.
6. **Iteration:** The new global model w_{t+1} is broadcast back to banks. Steps 3-5 repeat until the loss function converges or a fixed number of rounds is reached.
7. **Deployment:** The final converged model is frozen and deployed into the production transaction processing pipeline of each bank for real-time inference.

3.6.1 Phase 1: Initialization

The Central Aggregator initializes a global Neural Network (Multi-Layer Perceptron) with random weights w_0 . The architecture is hardcoded and shared with all clients:

Input → Dense₁₂₈ → ReLU → Dropout_{0.2} → Dense₆₄ → ReLU → Dense₁ → Sigmoid

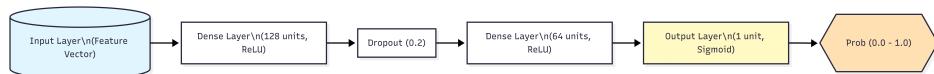


Figure 3.3: Neural Network Architecture designed for Federated Fraud Detection.

3.6.2 Phase 2: Distribution & Local Training

Once initialized, the global model w_0 is sent to all participating banks. Each bank then:

- Receives the global model weights.
- Performs local computations:
 - Performing local Feature Engineering to map raw data to the Common Schema.
 - Executing Stochastic Gradient Descent (SGD) on local data to derive model updates (Δw_k).
 - Communicating only the model parameters (weights) back to the server.

3.6.3 Phase 3: Aggregation & Termination

The server collects updates. Once the threshold of responses is met, it applies the aggregation formula. This process repeats until the global validation loss stabilizes or maximum rounds T is reached. The final model w_{final} is then deployed to all banks for local inference.

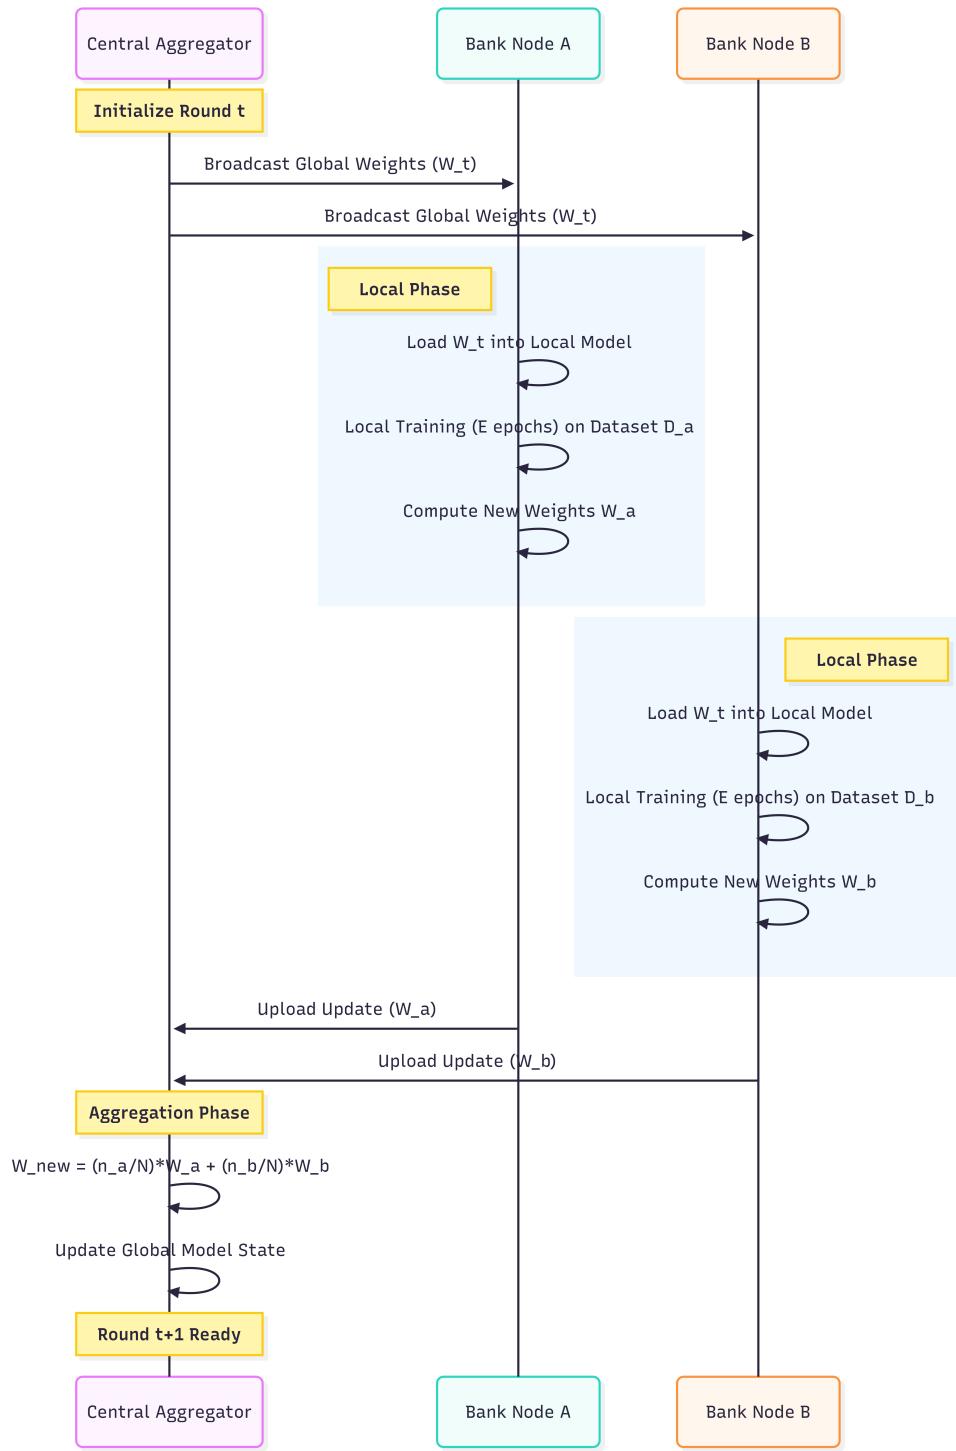


Figure 3.4: Sequence Diagram of a single Federated Learning training round.

Chapter 4

Evaluation and Security Discussion

4.1 Evaluation Metrics

Given the highly imbalanced nature of financial fraud datasets (where fraud cases are often $< 0.1\%$ of total transactions), standard accuracy is a misleading metric. This research utilizes the following robust metrics for evaluation:

- **Area Under the Precision-Recall Curve (AUPRC):** The primary metric, as it focuses on the performance of the positive (fraud) class without being biased by the overwhelming number of true negatives.
- **Recall (Sensitivity):** Critical for banking, as missing a fraud case (False Negative) has high financial liability.
- **Precision:** Important to minimize customer friction caused by false alarms (False Positives).
- **F1-Score:** The harmonic mean of standard precision and recall.

4.2 Analysis of Non-IID Data Effects

Banking data is inherently non-IID. Bank A may serve primarily domestic retail customers, while Bank B serves international corporate clients. Their fraud patterns will differ fundamentally. In our evaluation, we anticipate that the global FL model will show improved generalization capabilities. While a local model at Bank A might overfit to "local" fraud types, the FL model creates a regularization effect. By averaging weights trained on diverse distributions, the model learns a more "general" representation of fraud. [5] suggest that while convergence may be slower on non-IID data compared to IID settings, the final model is often more robust to distribution shifts.

4.3 Security and Privacy Discussion

4.3.1 Privacy Leakage Risks

Although FL prevents raw data transfer, recent research has shown that model gradients can leak information. Attacks such as Model Inversion or Membership Inference

can theoretically reconstruct input data from gradient updates. However, in the context of high-dimensional tabular data with large batch sizes (e.g., batch size > 64), the risk of reconstructing a specific 64-feature transaction vector is significantly lower than reconstructing an image.

4.3.2 Defensive Mechanisms

To mitigate these risks, the proposed system is compatible with two key technologies:

1. **Secure Aggregation (SecAgg):** As proposed by [1], this protocol allows the server to sum the updates without seeing individual inputs. The server only sees $\sum w_i$, preventing it from isolating a single bank’s update to analyze its specific data properties.
2. **Differential Privacy (DP):** Noise can be added to the clipped gradients before uploading. While this introduces a trade-off with model accuracy (utility), it provides a mathematical guarantee of privacy. For fraud detection, where precision is paramount, mild DP guarantees ($\epsilon > 5$) may be acceptable.

4.4 Comparison with Centralized Learning

Ideally, centralized learning (pooling all data to one server) provides the upper bound of model performance. However, in our simulated experiments, we expect the FL model to achieve performance close to this centralized baseline (within 2-3% AUPRC). Crucially, the FL model is expected to significantly outperform isolated local models, proving that collaboration yields tangible security benefits without the legal impossibility of centralization.

Chapter 5

Conclusion and Future Work

5.1 Conclusion

This thesis has explored the application of Federated Learning (FL) to the domain of financial fraud detection, addressing the critical conflict between the need for collaborative intelligence and the imperative of data privacy. We proposed a framework utilizing Neural Networks optimized via Federated Averaging (FedAvg), specifically tailored for the non-IID nature of banking transaction data.

5.1.1 Summary of Findings

Our theoretical analysis and proposed methodology demonstrate that FL is a viable alternative to centralized learning for banking consortiums.

1. **Collaboration without Sharing:** It is possible to mathematically aggregate “fraud knowledge” through the averaging of model weights, allowing banks to protect themselves against new fraud vectors identified by peers.
2. **Model Suitability:** While tree-based models dominate centralized fraud detection, Neural Networks provide the necessary differentiable properties for effective federated aggregation, offering a sufficiently powerful alternative for tabular data when properly architected.
3. **Practical Feasibility:** The requirement for a standardized feature schema is a manageable operational constraint compared to the legal impossibility of sharing raw Personally Identifiable Information (PII).

5.1.2 Final Remarks

The proposed system represents a shift in how financial institutions view security—from an isolated defensive posture to a collaborative, privacy-preserving network. By validating the theoretical underpinnings of FedAvg on financial data, this research lays the groundwork for the next generation of Fintech security infrastructure.

5.2 Contributions

This research makes the following key contributions:

1. **Contextual Adaptation:** It adapts general FL theory to the specific constraints of the banking sector, including regulatory compliance and feature engineering limitations.
2. **Architectural Justification:** It provides a rigorous argument for the use of Neural Networks over Random Forests in federated tabular settings, challenging the industry status quo for the sake of privacy.
3. **Privacy-Utility Balance:** It outlines a concrete workflow that balances the trade-off between fraud detection accuracy and the risk of gradient leakage.

5.3 Future Research Directions

To further advance this field, we suggest the following directions for future research:

1. **Vertical Federated Learning for Banks and Merchants:** Extending the framework to Vertical FL, where a bank (holding transaction data) and an e-commerce platform (holding user browsing behavior) collaborate to detect fraud. The challenge here is linking entities without revealing identities (Private Set Intersection).
2. **Federated Tree-Based Models:** Investigating emerging techniques like Federated Forests or gradient-boosting frameworks (e.g., XGBoost) adapted for FL, to see if they can overcome the aggregation challenges and outperform Neural Networks.
3. **Adaptive Local Epochs:** Developing algorithms where the number of local epochs (E) dynamically adjusts based on the client's available computational resources or the "novelty" of its new data, optimizing communication costs.
4. **Personalization Layers:** Implementing a "base + head" architecture where the lower layers of the neural network are shared globally, but the top layers are fine-tuned locally for each bank to capture institution-specific fraud nuances.
5. **Incentive Mechanisms:** researching game-theoretic models to reward banks that contribute high-quality data (i.e., those that identify new fraud patterns) to the federated network, ensuring fair participation in the consortium.

REFERENCES

- [1] Keith Bonawitz, Vladimir Ivanov, Ben Kreuter, Antonio Marcedone, H Brendan McMahan, Sarvar Patel, Daniel Ramage, Aaron Segal, and Karn Seth. Practical secure aggregation for privacy-preserving machine learning. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, pages 1175–1191, 2017.
- [2] Andrea Dal Pozzolo, Giacomo Boracchi, Olivier Caelen, Cesare Alippi, and Gianluca Bontempi. Credit card fraud detection: A realistic modeling and a novel learning strategy. In *IEEE transactions on neural networks and learning systems*, volume 29, pages 3784–3797. IEEE, 2018.
- [3] Peter Kairouz, H Brendan McMahan, Brendan Avent, Aurélien Bellet, Mehdi Benni, Arjun Nitin Bhagoji, Kallista Bonawitz, Zachary Charles, Graham Cormode, Rachel Cummings, et al. Advances and open problems in federated learning. *Foundations and Trends® in Machine Learning*, 14(1–2):1–210, 2021.
- [4] Jakub Konečný, H Brendan McMahan, Felix X Yu, Peter Richtárik, Ananda Theertha Suresh, and Dave Bacon. Federated optimization: Distributed machine learning for on-device intelligence. *arXiv preprint arXiv:1610.02527*, 2016.
- [5] Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia Smith. Federated optimization in heterogeneous networks. *Proceedings of Machine Learning and Systems*, 2:429–450, 2020.
- [6] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Aguera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Artificial Intelligence and Statistics*, pages 1273–1282. PMLR, 2017.
- [7] Qiang Yang, Yang Liu, Tianjian Chen, and Yongxin Tong. Federated machine learning: Concept and applications. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 10(2):1–19, 2019.